

Python Cheat Sheet for Data Science & Machine Learning

Pandas, Numpy, and Scikit-Learn are among the most popular libraries for data science and analysis with Python.

Numpy is used for lower level scientific computation. Pandas is built on top of Numpy and designed for practical data analysis in Python. Scikit-Learn comes with many machine learning models that you can use out of the box.

In this cheat sheet, we'll summarize some of the most common and useful functionality from these libraries. Let's jump straight in!

Setting Working Directory

```
import os
cwd = os.getcwd()
cwd
os.chdir('C:\\Users\\OneDrive\\Documents\\Data Science\\Datasets')
os.getcwd()
```

Importing Data

Any kind of data analysis starts with getting hold of some data. [Pandas](#) gives you plenty of options for getting data into your Python workbook:

```
import pandas as pd

df = pd.read_csv('SalaryData.csv')

df

df.shape
```

```
1 pd.read_csv(filename) # From a CSV file
2 pd.read_table(filename) # From a delimited text file (like TSV)
3 pd.read_excel(filename) # From an Excel file
4 pd.read_sql(query, connection_object) # Reads from a SQL table/database
5 pd.read_json(json_string) # Reads from a JSON formatted string, URL or file.
6 pd.read_html(url) # Parses an html URL, string or file and extracts tables to a list of dataframes
7 pd.read_clipboard() # Takes the contents of your clipboard and passes it to read_table()
8 pd.DataFrame(dict) # From a dict, keys for columns names, values for data as lists
```

Exploring Data

Once you have imported your data into a Pandas dataframe, you can use these methods to get a sense of what the data looks like:

```
1 df.shape() # Prints number of rows and columns in dataframe
2 df.head(n) # Prints first n rows of the DataFrame
3 df.tail(n) # Prints last n rows of the DataFrame
4 df.info() # Index, Datatype and Memory information
5 df.describe() # Summary statistics for numerical columns
6 s.value_counts(dropna=False) # Views unique values and counts
7 df.apply(pd.Series.value_counts) # Unique values and counts for all columns
8 df.describe() # Summary statistics for numerical columns
9 df.mean() # Returns the mean of all columns
10 df.corr() # Returns the correlation between columns in a DataFrame
11 df.count() # Returns the number of non-null values in each DataFrame column
12 df.max() # Returns the highest value in each column
13 df.min() # Returns the lowest value in each column
14 df.median() # Returns the median of each column
15 df.std() # Returns the standard deviation of each column
```

Selecting

Often, you might need to select a single element or a certain subset of the data to inspect it or perform further analysis. These methods will come in handy:

```
1 df[col] # Returns column with label col as Series
2 df[[col1, col2]] # Returns Columns as a new DataFrame
3 s.iloc[0] # Selection by position (selects first element)
4 s.loc[0] # Selection by index (selects element at index 0)
5 df.iloc[0,:] # First row
6 df.iloc[0,0] # First element of first column
```

Data Cleaning

If you're working with real world data, chances are you'll need to clean it up. These are some helpful methods:

```
1
2 df.columns = ['a','b','c'] # Renames columns
3 pd.isnull() # Checks for null Values, Returns Boolean Array
4 pd.notnull() # Opposite of s.isnull()
5 df.dropna() # Drops all rows that contain null values
6 df.dropna(axis=1) # Drops all columns that contain null values
7 df.dropna(axis=1,thresh=n) # Drops all rows have have less than n non null values
8 df.fillna(x) # Replaces all null values with x
9 s.fillna(s.mean()) # Replaces all null values with the mean (mean can be replaced with almost any function from the statistics
10 section)
11 s.astype(float) # Converts the datatype of the series to float
12 s.replace(1,'one') # Replaces all values equal to 1 with 'one'
13 s.replace([1,3],[ 'one','three']) # Replaces all 1 with 'one' and 3 with 'three'
14 df.rename(columns=lambda x: x + 1) # Mass renaming of columns
```

```
15 df.rename(columns={'old_name': 'new_name'}) # Selective renaming
    df.set_index('column_one') # Changes the index
    df.rename(index=lambda x: x + 1) # Mass renaming of index
```

Filter, Sort and Group By

Methods for filtering, sorting and grouping your data:

```
df[df[col] > 0.5] # Rows where the col column is greater than 0.5
1 df[(df[col] > 0.5) & (df[col] < 0.7)] # Rows where 0.5 < col < 0.7
2 df.sort_values(col1) # Sorts values by col1 in ascending order
3 df.sort_values(col2,ascending=False) # Sorts values by col2 in descending order
4 df.sort_values([col1,col2], ascending=[True,False]) # Sorts values by col1 in ascending order then col2 in descending order
5 df.groupby(col) # Returns a groupby object for values from one column
6 df.groupby([col1,col2]) # Returns a groupby object values from multiple columns
7 df.groupby(col1)[col2].mean() # Returns the mean of the values in col2, grouped by the values in col1 (mean can be replaced with
8 almost any function from the statistics section)
9 df.pivot_table(index=col1, values= col2,col3], aggfunc=mean) # Creates a pivot table that groups by col1 and calculates the mean
10 of col2 and col3
11 df.groupby(col1).agg(np.mean) # Finds the average across all columns for every unique column 1 group
12 df.apply(np.mean) # Applies a function across each column
    df.apply(np.max, axis=1) # Applies a function across each row
```

Joining and Combining

Methods for combining two dataframes:

```
1 df1.append(df2) # Adds the rows in df1 to the end of df2 (columns should be identical)
2 pd.concat([df1, df2],axis=1) # Adds the columns in df1 to the end of df2 (rows should be identical)
3 df1.join(df2,on=col1,how='inner') # SQL-style joins the columns in df1 with the columns on df2 where the rows for col have
    identical values. how can be one of 'left', 'right', 'outer', 'inner'
```

Writing Data

And finally, when you have produced results with your analysis, there are several ways you can export your data:

```
1 df.to_csv(filename) # Writes to a CSV file
2 df.to_excel(filename) # Writes to an Excel file
3 df.to_sql(table_name, connection_object) # Writes to a SQL table
4 df.to_json(filename) # Writes to a file in JSON format
5 df.to_html(filename) # Saves as an HTML table
6 df.to_clipboard() # Writes to the clipboard
```

Machine Learning

The Scikit-Learn library contains useful methods for training and applying machine learning models. For a complete list of the Supervised Learning, Unsupervised Learning, and Dataset Transformation, and Model Evaluation modules in Scikit-Learn, please refer to its [user guide](#).

```

1 # Import libraries and modules
2 import numpy as np
3 import pandas as pd
4
5 from sklearn.model_selection import train_test_split
6 from sklearn import preprocessing
7 from sklearn.ensemble import RandomForestRegressor
8 from sklearn.pipeline import make_pipeline
9 from sklearn.model_selection import GridSearchCV
10 from sklearn.metrics import mean_squared_error, r2_score
11 from sklearn.externals import joblib
12
13 # Load red wine data.
14 dataset_url = 'http://mlr.cs.umass.edu/ml/machine-learning-databases/wine-quality/winequality-red.csv'
15 data = pd.read_csv(dataset_url, sep=';')
16
17 # Split data into training and test sets
18 y = data.quality
19 X = data.drop('quality', axis=1)
20 X_train, X_test, y_train, y_test = train_test_split(X, y,
21                                                    test_size=0.2,
22                                                    random_state=123,
23                                                    stratify=y)
24
25 # Declare data preprocessing steps
26 pipeline = make_pipeline(preprocessing.StandardScaler(),
27                           RandomForestRegressor(n_estimators=100))
28
29 # Declare hyperparameters to tune
30 hyperparameters = { 'randomforestregressor__max_features': ['auto', 'sqrt', 'log2'],
31                    'randomforestregressor__max_depth': [None, 5, 3, 1]}
32
33 # Tune model using cross-validation pipeline
34 clf = GridSearchCV(pipeline, hyperparameters, cv=10)
35
36 clf.fit(X_train, y_train)
37
38 # Refit on the entire training set
39 # No additional code needed if clf.refit == True (default is True)
40
41 # Evaluate model pipeline on test data
42 pred = clf.predict(X_test)
43 print r2_score(y_test, pred)
44 print mean_squared_error(y_test, pred)
45
46 # Save model for future use
47 joblib.dump(clf, 'rf_regressor.pkl')
48 # To load: clf2 = joblib.load('rf_regressor.pkl')

```